

UG3 Introduction to Vision and Robotics

Vision Assignment

Clemens Wolff, Toms Bergmanis

March 7th, 2013

1 Introduction

The goal of this report is to propose an algorithm for object tracking. Object tracking, in general, is a challenging problem. Difficulties in tracking objects can arise due to abrupt object motion or changing appearance patterns of both the object and its surroundings and other factors. Tracking is usually performed in the context of higher-level applications that require the location and/or shape of the object in every frame. Typically, assumptions are made to constrain the tracking problem in the context of a particular application[1].

2 Methods

Several simplifying assumptions were made to constrain the tracking problem. It was assumed that:

1. objects to be tracked will be puck-like robots and that they will each appear in different shades of red or blue, or green.
2. camera will be set up in an angle not less than 45° with respect to the plane it is observing.
3. on the top of the robot will be a triangle indicating it's direction and that it will be in darker colour than the colour indicating the robot.
4. background of the image will be in colour other than any of the colours of the robots.

2.1 Detection of Robots

Input

I , a three channel image of dimensions $m \times n$ in the RGB colorspace.

Output

M , a $m \times n \times 3$ binary matrix where for each pixel P_{ij} of I , it holds that:

$M(i, j, 1) = 1 \leftrightarrow P_{ij}$ belongs to the red robot,

$M(i, j, 2) = 1 \leftrightarrow P_{ij}$ belongs to the green robot,

$M(i, j, 3) = 1 \leftrightarrow P_{ij}$ belongs to the blue robot.

Algorithm

1. Apply approximate RGB-normalisation to I , giving I_n :
 - For each pixel in I_n , calculate the sum S_{rgb} of the red, green, and blue values of that pixel.
 - If $S_{rgb} \neq 0$ (the pixel is not absolute black), set each of the pixel's red, green, and blue values to that value divided by S_{rgb} .
2. Calculate μ_r, μ_g, μ_b and $\sigma_r, \sigma_g, \sigma_b$, the means and standard deviations of the values in the three channels of I_n .
3. Assign each pixel P_{ij} in I to one of the robots or to the background:
 - Normalise P 's red, green, and blue values, giving P_n .
 - Calculate the probabilities p_r, p_g, p_b that P_n was generated by the gaussian distributions $\mathcal{N}_r = (\mu_r, \sigma_r), \mathcal{N}_g = (\mu_g, \sigma_g), \mathcal{N}_b = (\mu_b, \sigma_b)$.
 - Calculate P 's hue value h .
 - If h is within a certain range defined as red and p_r is sufficiently small, set $M(i, j, 1) = 1$ (similarly for ranges defined as green/blue and p_g/p_b). If none of these conditions are met, set $M(i, j, 1) = M(i, j, 2) = M(i, j, 3) = 0$.
4. Remove noise from each channel in M :
 - Set pixels to zero if they have fewer neighbours with value one than they have adjacent pixels with value zero.
 - Set zero-valued pixels to one if they have two one-valued horizontal or vertical neighbours.
5. Remove components that are distant from the main concentration of mass in each channel in M :
 - Compute the center of mass C of the channel.
 - Compute c_1, c_2, \dots , the centers of mass of each connected component in the channel.
 - Compute the mean distance δ of the c_k to C .
 - Set $M(i, j) = 0$ for all the pixels (i, j) in the components k that satisfy $c_k > \tau\delta$ for some fixed threshold τ .
6. Exploit the fact that all robots have similar sizes by setting every channel in M to all-zeros if the number of pixels set in that channel is smaller than the number of pixels set in the most populated channel by some margin.

Figure 1 shows a visualisation of the output matrix M .

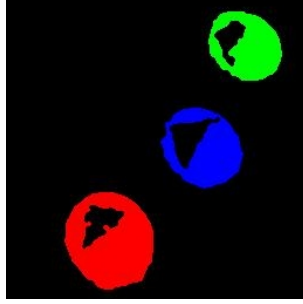


Figure 1: Result of Robot detection

2.2 Finding Robot Directions

Input

I , a three channel image of dimensions $m \times n$ in the RGB colorspace.

Output

$\Lambda = \{(c_r^m, c_r^t), (c_g^m, c_g^t), (c_b^m, c_b^t)\}$, a set where c_r^m is the center of mass of the red robot and c_r^t is the point towards which the robot is facing (similarly for the green and blue robots).

Algorithm

1. Get a matrix of robot masks M using the algorithm in Section 2.1. Let M_i be the i^{th} channel of M i.e. the set of points $\{M(a, b, i) | 1 \leq a \leq m, 1 \leq b \leq n\}$. Apply the remainder of the algorithm to each channel ξ in M .
2. Calculate the convex hull H of the points in the channel and create the set of pixels of I that are inside H : $P = \{p_{ij} | p_{ij} \in \xi \wedge M(i, j, \xi) = 1\}$
3. Calculate μ , the average rgb-value over P . Generate $\Pi = \{p | p \in P \wedge rgbvalue(p) < \mu\}$, the set of pixels in P that have a below-average rgb value.
4. The black triangles on the robots are the pixels in Π . Get rid of them by setting the relevant indices in M to zero. Recompute the convex hull of M .
5. Repeat the previous step and remember the pixels in Π .
This reduces noise in M by giving a tighter estimate on the robot's pixels when the triangles were under-detected by the algorithm in Section 2.1.
Figure 2 shows the result of this step - a notable improvement in clarity of the triangles compared to Figure 1.
6. Update Λ : c_ξ^m is the center of mass of M_ξ , c_ξ^t is the center of mass of Π .
A line from c_ξ^m to c_ξ^t indicates the direction of the robot.

Figure 3 shows a visualisation of the output set Λ .

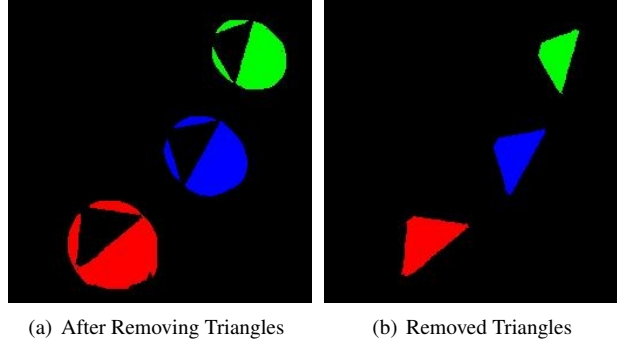


Figure 2: Triangle Detection via Local Thresholding




Figure 3: Detected Directions

2.3 Tracking Robots Over a Sequences of Frames

Input

$\Upsilon = \{I_1, I_2, \dots\}$, a sequence where each of the I_i is a three channel image of dimensions $m \times n$ in the RGB colorspace.

Output

Ω , a visualisation of the robot positions over Υ .

Algorithm

1. Use a median-filter to generate a background Ω from Υ .
 For each $1 \leq i \leq m, 1 \leq j \leq n$:
 - Create $\omega_{ij} = \{I_k(i, j) | I_k \in \Upsilon\}$, the set of the colors of the pixels at location (i, j) of all the images in Υ .
 - Set $\Omega(i, j) = \text{median}(\omega_{ij})$.
2. For each $I_i \in \Upsilon$:
 - Use the algorithm in Section 2.2 to get the set Λ . Let $\lambda = \{c | (c, -) \in \Lambda\}$.
 - Overlay Ω with a line from each element in λ_{i-1} to the corresponding element in λ_i , thus linking the centroids from image I_{i-1} to the centroids in image I_i .

Figure 4 shows a visualisation of the resulting track.

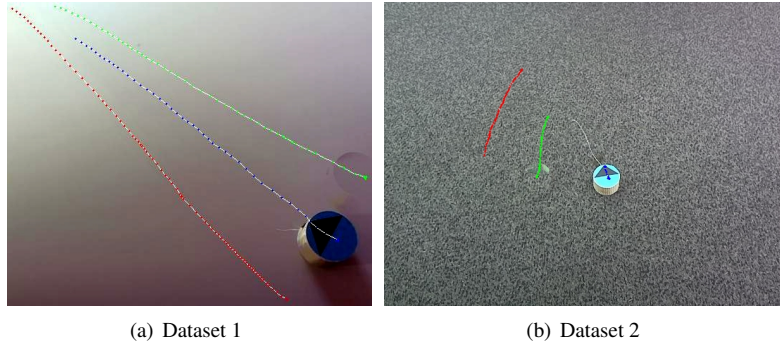


Figure 4: Output of Tracing Algorithm

3 Results

This section evaluates and visualises the performance of the three algorithms presented in Sections 2.1, 2.2, and 2.3. Table 1 describes the properties of the datasets used for this evaluation.

#	Background	Robot Size	Robot Color	Illumination
1	uniform, gray	large	saturated, dark	uniform, red hue
2	noisy, gray	small	faded, blue robot is cyan	histograms are bell-shaped
3	patterned, brown	large	saturated	daylight only
4	patterned, brown	large	saturated	daylight and artificial light

Table 1: Properties of evaluation datasets

3.1 Detection of Robots

The algorithm described in Section 2.1, worked perfectly on datasets 1 and 2. Evaluation on the third dataset led to the worst performance over all datasets, with $\sim 60\%$ of the occurrences of the blue robot being undetected and $\sim 40\%$ of the occurrences of the green robot being under-detected (leading to bad direction detection). The performance on the fourth dataset was interesting: the red robot was under-detected in $\sim 45\%$ of the cases (with the blue and green robots being found just fine) - while in the other datasets the red robot was usually detected with the highest confidence. Over all four datasets, about 10% of the robot instances were badly detected.

The fact that the color-detection algorithm works well on both datasets 1 and 2 leads to the conjecture that it is invariant under texture changes in the scene background and variations in robot-color saturation. The bad performance on dataset 3 can be explained by interference from the color of the scene background and by changes in scene illumination. The fact that the algorithm offers almost top-level performance on dataset 4 (captured on the same background as dataset 3) implies that the change in scene illumination is probably the largest influence on the algorithm’s performance.

This is in keeping with the intuition that daylight has more inherent variation than arti-

ficial light, thus introducing a higher degree of variability into the characteristics of the captured images.

3.2 Detection of Directions

Performance of robot direction detection, understandably, is heavily dependent on the performance of the detection of the robots. If the robots are well isolated by Section 2.1's algorithm, robot orientations are perfectly detected.

The algorithm is invariant under loose detection - false positive cases where some addition non-robot region is misleadingly detected as a robot. This is due to the algorithm's ability to filter-out noisy detections.

In case of under-detection - false negative cases where some parts of the image representing the robot where not detected - the algorithm breaks: the predicted direction is skewed towards the opposite side of the under-detection. This is due to the algorithm not putting strict circular or elipsoidal constraints on the shape of the robots: under detection is thus able to move the center of mass of either the triangle or the robot-convex-hull. The error introduced by this is proportional to the area of the under-detected region.

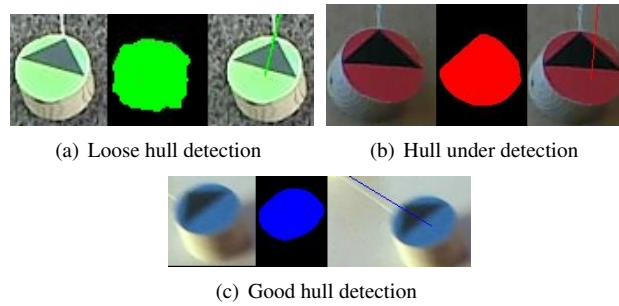


Figure 5: Direction detections for convex hulls of different qualities

3.3 Tracking of the robots

Section 2.3's algorithm to track robots over consecutive frames is trivial - a mere visualisation of half of the results of the robot-direction- detection algorithm presented in Section 2.2. The tracking algorithm's performance is therefore directly related to the performance of the robot-direction-detection algorithm and the same observations as in Section 3.2 apply: generally speaking, the algorithm performed well.

Figure 4 begets one additional observation related to the evaluation of the robot-tracking algorithm: both datasets considered in this report exhibit the property that one of the objects of interest does not move much for most of the frames. This entails that generating a background from data employing a simple frame-difference base approach (such as the median- filter used in Section 2.3) to perform background subtraction is bound to fail as one of the objects of interest will be considered a part of the background due to being mostly stationary. This is unfortunate since pre-processing the datasets with background subtraction would increase the accuracy of Section 2.1's algorithm by reducing noise and increasing resolution in the image.

4 Discussion

The algorithms in Sections 2.1, 2.2, and 2.3 operate under a limited number of reasonable simplifying assumptions and performed well across a range of datasets captured under very different conditions.

The robot color detection algorithm of Section 2.1 could be improved by adding more stringent conditions on the shape of the robots e.g. fitting a circle or ellipse to them rather than a convex hull. This should improve the subsequent performance of the direction-detection algorithm by reducing the amount of under-detections.

If assumptions about the size, scale, and shape of the robots can be made, Hough transforms or other shape-based techniques could be used to detect the robots. This would be more robust than the current color/pixel based approach but not invariant under scale or differences in camera positions.

Another way to improve the detection of the robots could be augmenting the current approach with the utilisation of second order spatial image statistics to pre-process the images. This would lead to a rough estimate of the robot positions with few false negatives, thus improving the currently utilised algorithms due to a reduced search space (implying a denser signal). A reduced search space also allows for computationally expensive but high-precision techniques to be used. One example of such a technique is a local color-based search starting from a small seed of pixels that are hypothesised to be part of a robot.

Direction detection could be improved by making it less dependent on the accuracy of the robot detection. One way to achieve this would be to base the direction calculation on properties of the detected triangles (e.g. direction = tangent line to the longest side) rather than on the triangles' centers of mass. This would increase robustness to under-detections.

Cross-frame tracking could be improved by abusing spatio-temporal proximity. One way to do this would be put a cap on how far the centroid of the robot can move in a certain frame interval, thus smoothing out the odd completely wrong prediction.

References

- [1] Yilmaz, A., Javed, O., and Shah, M., "Object Tracking: A Survey", *ACM Comput. Surv.* 38, 4, Article 13, 2006.

```

1 function res = main(path, image_type, start_offset, time_step)
2     if nargin < 2
3         image_type = 'jpg';
4     end
5     if nargin < 3
6         start_offset = 0;
7     end
8     if nargin < 4
9         time_step = 0.33;
10    end
11
12    files = dir(sprintf('%s/*.%s', path, image_type));
13    filenames = {files.name};
14    [~, num_files] = size(filenames);
15
16    [m n ~] = size(imread(sprintf('%s/%s', path, filenames{1})));
17    track_mask = zeros(m, n);
18
19    rc = zeros(num_files - start_offset, 2);
20    gc = zeros(num_files - start_offset, 2);
21    bc = zeros(num_files - start_offset, 2);
22
23    reds = cell(num_files - start_offset, 1);
24    greens = cell(num_files - start_offset, 1);
25    blues = cell(num_files - start_offset, 1);
26
27    for i = 1 + start_offset : num_files
28        image = imread(sprintf('%s/%s', path, filenames{i}));
29        [directions, centroids] = analyse_image(image);
30
31        r = uint16(centroids(1,:));
32        if r(1) > 0 && r(2) > 0
33            track_mask = overlay_cross(track_mask, 1, r(1), r(2));
34            rc(i,:) = [r(1), r(2)];
35        end
36
37        g = uint16(centroids(2,:));
38        if g(1) > 0 && g(2) > 0
39            track_mask = overlay_cross(track_mask, 2, g(1), g(2));
40            gc(i,:) = [g(1), g(2)];
41        end
42
43        b = uint16(centroids(3,:));
44        if b(1) > 0 && b(2) > 0
45            track_mask = overlay_cross(track_mask, 3, b(1), b(2));
46            bc(i,:) = [b(1), b(2)];
47        end
48
49        reds{i} = image(:,:,1);
50        greens{i} = image(:,:,2);
51        blues{i} = image(:,:,3);
52        imshow(overlay_mask(image, directions));
53
54        pause(time_step);
55    end
56
57    rc = rc(any(rc,2),:);
58    gc = gc(any(gc,2),:);
59    bc = bc(any(bc,2),:);
60
61    red_median = median(cat(3, reds{:}), 3);
62    green_median = median(cat(3, greens{:}), 3);
63    blue_median = median(cat(3, blues{:}), 3);
64

```



```

65     bg = cat(3, red_median, green_median, blue_median);
66
67     bg = overlay_polygon(bg, rc, [255, 255, 255]);
68     bg = overlay_polygon(bg, gc, [255, 255, 255]);
69     bg = overlay_polygon(bg, bc, [255, 255, 255]);
70     bg = overlay_mask(bg, track_mask);
71     imshow(bg)
72 end
73
74
75 function [pretty_mask, varargout] = analyse_image(image)
76     [num_rows, num_cols, num_channels] = size(image);
77     blob_mask = mask_colors(image);
78     [convex_mask, ~] = mask_convex_regions(image, blob_mask);
79     [convex_mask] = demask_triangles(image, convex_mask, false);
80     [convex_mask, ~] = mask_convex_regions(image, convex_mask);
81     [~, centroids, ~, triangle_centroids] = ...
82         demask_triangles(image, convex_mask, true);
83     pretty_mask = overlay_rays(zeros(num_rows, num_cols, num_channels), ...
84         centroids, triangle_centroids, 99, ...
85         'Color', [1 0 0; 0 1 0; 0 0 1]);
86     varargout{1} = centroids;
87     varargout{2} = triangle_centroids;
88     varargout{3} = convex_mask;
89 end
90
91
92 % this function can used to find both - triangles and masks without
93 % triangles. To find triangles it must be run with
94 % get_triangle_centroids=true. Also application of the function int the
95 % following fashion gives better results.
96 %     [convex_mask, ~] = mask_convex_regions(image, blob_mask);
97 %     [convex_mask] = demask_triangles(image, convex_mask, false);
98 %     [convex_mask, ~] = mask_convex_regions(image, convex_mask);
99 function [mask, varargout] = demask_triangles(image, mask, ...
100         get_triangle_centroids)
101     [num_rows, num_cols, num_channels] = size(image);
102     centroids = zeros(num_channels, 2);
103     triangle_centroids = zeros(num_channels, 2);
104     trinagle_mask = mask;
105     for c = 1 : num_channels
106         channel = image(:,:,c);
107         channel_mask = mask(:,:,c);
108         channel_trinagle_mask = mask(:,:,c);
109         if ~any(channel_mask)
110             continue;
111         end
112         rgb_values = zeros(sum(channel_mask(:)), 1);
113         idx = 1;
114         for nr = 1 : num_rows
115             for nc = 1 : num_cols
116                 if channel_mask(nr, nc) == 0
117                     continue;
118                 end
119                 rgb_values(idx) = channel(nr, nc);
120                 idx = idx + 1;
121             end
122         end
123         mean_rgb = mean(rgb_values(:));
124         channel_mask(channel < mean_rgb) = 0;
125         mask(:,:,c) = channel_mask;
126         props = regionprops(channel_mask, 'Centroid');
127         centroid = props.Centroid;
128         centroids(c,:) = [centroid(2), centroid(1)];
129         if get_triangle_centroids

```

```

130         channel_trinagle_mask(channel > mean_rgb) = 0;
131         channel_trinagle_mask=filter_mask(channel_trinagle_mask);
132         trinagle_mask(:, :, c) = channel_trinagle_mask;
133         props = regionprops(channel_trinagle_mask, 'Centroid');
134         centroid = props.Centroid;
135         triangle_centroids(c, :) = [centroid(2), centroid(1)];
136     end
137 end
138 varargout{1} = centroids;
139 varargout{2} = trinagle_mask;
140 varargout{3} = triangle_centroids;
141 end
142
143
144 function [color_mask, varargout] = mask_convex_regions(image, mask)
145     [num_rows, num_cols, num_channels] = size(image);
146     color_mask = zeros(num_rows, num_cols, num_channels);
147     convex_centroids = zeros(num_channels, 2);
148     for c = 1 : num_channels
149         channel = mask(:, :, c);
150         if ~any(channel(:))
151             continue;
152         end
153         props = regionprops(channel, 'Centroid', 'ConvexImage', 'BoundingBox');
154         convex_props = regionprops(props.ConvexImage, 'Centroid');
155         convex_centroid = convex_props.Centroid;
156         convex_centroid = [convex_centroid(2) + props.BoundingBox(2), ...
157             convex_centroid(1) + props.BoundingBox(1)];
158         convex_centroids(c, :) = convex_centroid;
159         convex_image = props.ConvexImage;
160         [num_rows_convex, num_cols_convex] = size(convex_image);
161         for row = 1 : num_rows_convex
162             for col = 1 : num_cols_convex
163                 if convex_image(row, col) == 1
164                     newrow = round(row + props.BoundingBox(2));
165                     newcol = round(col + props.BoundingBox(1));
166                     color_mask(newrow, newcol, c) = 1;
167                 end
168             end
169         end
170     end
171     varargout{1} = convex_centroids;
172 end
173
174
175 function image_mask = mask_colors(image)
176     [num_rows, num_cols, ~] = size(image);
177     num_pixels = num_rows * num_cols;
178     image_mask = zeros(num_pixels, 3);
179     rgb = double(reshape(image, num_pixels, 3));
180     rgbN = double(reshape(normalise_rgb(image, 'approximate'), num_pixels, 3));
181     rN_sdev = std(rgbN(:, 1));
182     gN_sdev = std(rgbN(:, 2));
183     bN_sdev = std(rgbN(:, 3));
184     rN_mean = mean(rgbN(:, 1));
185     gN_mean = mean(rgbN(:, 2));
186     bN_mean = mean(rgbN(:, 3));
187     hsv = reshape(rgb2hsv(image), num_pixels, 3);
188     for c = 1 : num_pixels
189         rN = rgbN(c, 1);
190         gN = rgbN(c, 2);
191         bN = rgbN(c, 3);
192         hue = hsv(c, 1) * 360;
193         % current pixel is red
194         if (hue >= 330 || hue <= 30) && ...

```

```

195         (normal_prob(rN, rN_mean, rN_sdev) < 0.001)
196         image_mask(c,1) = 1;
197     % current pixel is green
198     elseif (hue >= 80 && hue < 180) && ...
199         (normal_prob(gN, gN_mean, gN_sdev) < 0.007)
200         image_mask(c,2) = 1;
201     % current pixel is blue
202     elseif (hue >= 150 && hue <= 270) && ...
203         (normal_prob(bN, bN_mean, bN_sdev) < 0.0000085)
204         image_mask(c,3) = 1;
205     end
206 end
207 image_mask = reshape(image_mask, num_rows, num_cols, 3);
208 image_mask = remove_noise(image_mask);
209 image_mask = remove_outliers(image_mask);
210 image_mask = enforce_similar_channel_areas(image_mask);
211 end
212
213 function x = normal_prob(val, mu, sigma)
214     x = 1.0 / (sigma * sqrt(2 * pi)) * exp(-(val - mu) ^ 2 / (2 * sigma ^ 2));
215 end
216
217
218 function image = remove_noise(image)
219     [~, ~, num_channels] = size(image);
220     for c = 1 : num_channels
221         channel = image(:,:,c);
222         channel = bwmorph(channel, 'majority', Inf);
223         channel = bwmorph(channel, 'bridge', Inf);
224         image(:,:,c) = channel;
225     end
226 end
227
228
229 % finds the connected components in each channel of |image| and removes those
230 % that are far away from the centroid of the pixels in that channel
231 % here, 'far away' means more distant than |distance_proportion_threshold| times
232 % the average distance of each connected component to the channel centroid
233 % this removes big areas of noise such as the big green blob inside of the black
234 % arrow of the red robot in data/1/00000006.jpg
235 function image = remove_outliers(image, distance_proportion_threshold)
236     if nargin < 2
237         distance_proportion_threshold = 0.5;
238     end
239     [~, ~, num_channels] = size(image);
240     for c = 1 : num_channels
241         channel = image(:,:,c);
242         if ~any(channel(:))
243             continue;
244         end
245         channel_properties = regionprops(channel, 'Centroid');
246         channel_centroid = channel_properties.Centroid;
247         regions = bwconncomp(channel);
248         regions_properties = regionprops(regions, 'Centroid', 'PixelIdxList');
249         regions_centroids = {regions_properties.Centroid};
250         distances = cellfun(@(x) norm(x - channel_centroid), regions_centroids);
251         mean_distance = mean(distances);
252         for d = 1 : length(distances)
253             if distances(d) > mean_distance * distance_proportion_threshold
254                 idx = regions_properties(d).PixelIdxList;
255                 channel(idx) = 0;
256             end
257         end
258     end
259     image(:,:,c) = channel;

```

```

260     end
261 end
262
263
264 % we know that the robots are all about the same size - we can thus remove any
265 % channels in the mask that have a much smaller area than the other channels
266 % this catches some problems like the shadow of the blue robot in
267 % data/1/00000095.jpg being detected as a red blob
268 function image = enforce_similar_channel_areas(image, area_proportion_threshold)
269     if nargin < 2
270         area_proportion_threshold = 0.5;
271     end
272     [~, ~, num_channels] = size(image);
273     areas = zeros(num_channels, 1);
274     for c = 1 : num_channels
275         channel = image(:,:,c);
276         if ~any(channel(:))
277             continue;
278         end
279         region_props = regionprops(channel, 'Area');
280         areas(c) = region_props.Area;
281     end
282     max_area = max(areas(areas > 0));
283     for c = 1 : num_channels
284         area = areas(c);
285         if area == 0
286             continue;
287         end
288         if (area < max_area * area_proportion_threshold) || ...
289             (area > max_area / area_proportion_threshold)
290             image(:,:,c) = image(:,:,c) * 0;
291         end
292     end
293 end
294
295
296 % this filters out all connected regions but the biggest one
297 function mask = filter_mask(mask)
298     [x, y, num_channels] = size(mask);
299     for c = 1 : num_channels
300         channel = zeros(x, y);
301         channel = reshape(channel, x * y, 1);
302         blob_info = bwconncomp(mask(:,:,c));
303         blob_list = blob_info.PixelIdxList;
304         [nrows, ncols] = cellfun(@size, blob_list);
305         largest_blob_pixels = blob_list{find(nrows == max(nrows))};
306         for i = 1 : max(nrows)
307             channel(largest_blob_pixels(i)) = 1;
308         end
309         channel = reshape(channel, x, y);
310         mask(:,:,c) = channel;
311     end
312 end
313
314
315 % normalises the values of the red, green, and blue channels of |image| in order
316 % to eliminate illumination differences in the image
317 % formula used: {r, g, b} = {r, g, b} / sqrt(r^2 + g^2 + b^2)
318 % if 'approximate' is passed as an additional parameter, instead use
319 % {r, g, b} = {r, g, b} / (r + g + b)
320 % this is approximately two times faster than the exact normalisation
321 function normalised_image = normalise_rgb(image, varargin)
322     approximate = ~isempty(find(strcmpi(varargin, 'approximate')));
323     red = double(image(:,:,1));
324     green = double(image(:,:,2));

```

```

325 blue = double(image(:,:,3));
326 if approximate
327     euclid_rgb = red(:,:,) + green(:,:,) + blue(:,:,);
328 else
329     euclid_rgb = sqrt(red(:,:,).^2 + green(:,:,).^2 + blue(:,:,).^2);
330 end
331 red_norm = round(red(:,:,) ./ euclid_rgb .* 255);
332 green_norm = round(green(:,:,) ./ euclid_rgb .* 255);
333 blue_norm = round(blue(:,:,) ./ euclid_rgb .* 255);
334 % some pixels are absolute black (r = g = b = 0) which causes division by
335 % zero errors during normalisation and NaN values in the normalised channels
336 % need to filter these values out
337 red_norm(isnan(red_norm)) = 0;
338 green_norm(isnan(green_norm)) = 0;
339 blue_norm(isnan(blue_norm)) = 0;
340 red_norm = uint8(red_norm);
341 green_norm = uint8(green_norm);
342 blue_norm = uint8(blue_norm);
343 normalised_image = cat(3, red_norm, green_norm, blue_norm);
344 end
345
346 % returns image with a cross centered on pixel (|x|, |y|) drawn in |channel|
347 % [0, 0] is the top left corner of the image
348 function image = overlay_cross(image, channel, x, y)
349     [h w ~] = size(image);
350
351     if channel == 1
352         other1 = 2;
353         other2 = 3;
354     end
355     if channel == 2
356         other1 = 1;
357         other2 = 3;
358     end
359     if channel == 3
360         other1 = 1;
361         other2 = 2;
362     end
363
364     if y + 1 < h
365         image(x, y + 1, channel) = 1;
366         image(x, y + 1, other1) = 0;
367         image(x, y + 1, other2) = 0;
368     end
369
370     if y - 1 > 0
371         image(x, y - 1, channel) = 1;
372         image(x, y - 1, other1) = 0;
373         image(x, y - 1, other2) = 0;
374     end
375
376     if x + 1 < w
377         image(x + 1, y, channel) = 1;
378         image(x + 1, y, other1) = 0;
379         image(x + 1, y, other2) = 0;
380     end
381
382     if x - 1 > 0
383         image(x - 1, y, channel) = 1;
384         image(x - 1, y, other1) = 0;
385         image(x - 1, y, other2) = 0;
386     end
387
388     image(x, y, channel) = 1;

```

```

390     image(x, y, other1) = 0;
391     image(x, y, other2) = 0;
392
393 end
394
395
396 % returns the result of putting |mask| onto |image|
397 % for each pixel that is set in some channel of |mask|, saturates the pixel in
398 % the equivalent channel of |image|
399 function image = overlay_mask(image, mask, varargin)
400     % parse options
401     argc = size(varargin, 2);
402     c = 1;
403     while c <= argc
404         arg = varargin{c};
405         if strcmpi(arg, 'GrayScale')
406             grayscale = 1;
407         elseif strcmpi(arg, 'Saturation')
408             if c + 1 > argc
409                 error('Saturation option should be followed by a double');
410             end
411             saturation = varargin{c + 1};
412             c = c + 1;
413         elseif strcmpi(arg, 'Lightness')
414             if c + 1 > argc
415                 error('Lightness option should be followed by a double');
416             end
417             lightness = varargin{c + 1};
418             c = c + 1;
419         end
420         c = c + 1;
421     end
422
423     % modify background image
424     if exist('grayscale', 'var')
425         gray_image = rgb2gray(image);
426         image = cat(3, gray_image, gray_image, gray_image);
427     end
428     if exist('saturation', 'var')
429         hsv_image = rgb2hsv(image);
430         hsv_image(:,:,2) = hsv_image(:,:,2) * saturation;
431         image = hsv2rgb(hsv_image);
432     end
433     if exist('lightness', 'var')
434         hsv_image = rgb2hsv(image);
435         hsv_image(:,:,3) = hsv_image(:,:,3) * lightness;
436         image = hsv2rgb(hsv_image);
437     end
438
439     % lay mask onto background image
440     num_channels = size(image, 3);
441     channels = 1 : num_channels;
442     for c = 1 : num_channels
443         channel = image(:,:,c);
444         mask_pixels = find(mask(:,:,c) == 1);
445         channel(mask_pixels) = 255;
446         image(:,:,c) = channel;
447         for d = setdiff(channels, c)
448             channel = image(:,:,d);
449             channel(mask_pixels) = 0;
450             image(:,:,d) = channel;
451         end
452     end
453 end
454

```

```

455
456 % returns image with lines drawn from the nth point in |points| to the n+1th
457 % [0, 0] is the top left corner of the image
458 function image = overlay_polygon(image, points, color)
459     if nargin < 3
460         color = [255 255 255];
461     end
462
463     points = round(points);
464
465     num_channels = size(image, 3);
466     for c = 1 : num_channels
467         channel = image(:,:,c);
468         channel = overlay_polygon_channel(channel, points, color(c));
469         image(:,:,c) = channel;
470     end
471 end
472
473
474 % Bresenham's line algorithm (simplified version)
475 % http://en.wikipedia.org/wiki/Bresenham's\_line\_algorithm#Simplification
476 function channel = overlay_polygon_channel(channel, points, color)
477     [xmax, ymax] = size(channel);
478     for c = 1 : length(points) - 1
479         start = points(c,:);
480         stop = points(c + 1,:);
481         x0 = start(1);
482         y0 = start(2);
483         x1 = stop(1);
484         y1 = stop(2);
485
486         dx = abs(x1 - x0);
487         dy = abs(y1 - y0);
488
489         if x0 < x1
490             sx = 1;
491         else
492             sx = -1;
493         end
494         if y0 < y1
495             sy = 1;
496         else
497             sy = -1;
498         end
499
500         err = dx - dy;
501
502         while 1
503             if x0 > 0 && x0 <= xmax && y0 > 0 && y0 <= ymax
504                 channel(x0, y0) = color;
505             end
506             if x0 == x1 && y0 == y1
507                 break;
508             end
509             e2 = 2 * err;
510             if e2 > -dy
511                 err = err - dy;
512                 x0 = x0 + sx;
513             end
514             if e2 < dx
515                 err = err + dx;
516                 y0 = y0 + sy;
517             end
518         end
519     end

```

```

520 end
521
522
523 function image = overlay_rays(image, from, to, length, varargin)
524     % parse options
525     argc = size(varargin, 2);
526     c = 1;
527     while c <= argc
528         arg = varargin{c};
529         if strcmpi(arg, 'Color')
530             if c + 1 > argc
531                 error('Color option should be followed by an integer tripplet');
532             end
533             colors = varargin{c + 1};
534             c = c + 1;
535         end
536         c = c + 1;
537     end
538     % option defaults
539     if ~exist('colors', 'var')
540         colors = [255 255 255];
541     end
542
543     num_rays = size(from, 1);
544     if size(colors, 1) == 1;
545         colors = repmat(colors, num_rays, 1);
546     end
547
548     for c = 1 : num_rays
549         if from(c,:) == to(c,:)
550             continue;
551         end
552         x0 = from(c,1);
553         y0 = from(c,2);
554         x1 = to(c,1);
555         y1 = to(c,2);
556         dx = x0 - x1;
557         dy = y0 - y1;
558         lambda = min(sqrt(length ^ 2 / (dx ^ 2 + dy ^ 2)), ...
559                     -sqrt(length ^ 2 / (dx ^ 2 + dy ^ 2)));
560
561         x = x0 + lambda * dx;
562         y = y0 + lambda * dy;
563
564         image = overlay_polygon(image, [x0 y0; x y], colors(c,:));
565     end
566 end

```